

AI Work & Vision

An origin story, the framework, the platform — and what I believe about the next 36 months.

PREPARED

April 2026

I've spent the last 18 months building a **personal operating system** on top of AI. Out of that came a methodology for how operators should deploy AI inside their organizations — and a platform venture I'm building separately. The methodology travels with any technology. The platform is **mception**.

— ORIGIN STORY

How this actually got built — seven steps over 18 months.

01

I wired frameworks into my OS — product innovation, strategy, ops methods from classes, books, and 16 years of operating.

Not theory in a notebook. Actual prompts, agents, and decision rules running on top of my daily work, accumulating context every day.

FOUNDATION • 2024-2025

02

A friend leaving CPG asked what was next. I asked my OS to generate ideas using my frameworks.

Not a search query. The OS pulled from the trend layer, ran the ideation methods I'd loaded, and produced a handful of actually interesting concept descriptions — the kind a workshop would have taken three weeks to generate.

TRIGGER • Q4 2025

03

The output was good enough that I deepened it: more frameworks, more research, image generation, a polished PDF deliverable.

Each upgrade was a few hours of work, not weeks. The system got better every cycle because the frameworks accumulated and the research layer learned what mattered.

COMPOUNDING • EARLY 2026

04

I wrapped it in a web tool and kept iterating in small daily increments.

Recursive learning kicked in: one run informed the next. The deliverable quality climbed without me rebuilding anything.

PRODUCTIZATION • FEBRUARY 2026

05

One conversation with Panda Restaurant Group — three hours later I had replicated the engine for QSR operations innovation.

Same architecture, different problem space. Product innovation logic became operations innovation logic with a different operator-wisdom library and a different scoring rubric.

Three hours, not three weeks.

REPLICATION • APRIL 2026

06

A furniture operator reached out. The product and ops engines blended naturally into one engagement.

Innovation workshop output — 10 concepts, hero imagery, full pitch decks — landed inside a week. The frameworks didn't need to be rebuilt. They composed.

COMPOSITION • APRIL 2026

07

Today: ~80% automated end-to-end. The remaining 20% is where I want my judgment in the loop.

The system goes broader and deeper on research, saves it, triggers internal debate or human feedback, and incorporates the feedback in the next cycle — in hours, not weeks.

STATUS • ONGOING

— VISUAL SAMPLE

What the engine produces — drawn from a recent retail engagement.



CONCEPT • CONCIERGE
APPOINTMENT SHOWROOM



CONCEPT • 15-MINUTE SMALL-
FORMAT SHOWROOM



CONCEPT • SUBSCRIPTION /
RENT-TO-OWN TIER

Three of ten concepts generated end-to-end — research, framework match, hero imagery, RTBs, executive summary, whitepaper appendix. Total cycle: under one week.

— THE ENGINE

A recursive learning loop — already 80% automated.



The compounding effect. Every cycle leaves the system smarter. The frameworks library grows. The research library grows. The agent library grows. What used to take a workshop now takes a session — and the next session is better than the last because the system remembers.

— TWO THINGS, NOT ONE

The framework is the methodology. The platform is the venture.

THE FRAMEWORK

Methodology

A way of thinking about how decisions get made when AI is embedded in the operating model — what agents own, where humans are irreplaceable, how institutional memory accumulates, and how the system measures whether it's making you better over time.

THE PLATFORM

mception

A separate LLC. Independent venture. Self-funded. Infrastructure that operationalizes the framework into deployable form so other operators can eventually run it.

Small group of advisors. *No formal partners, no employees, no vendors yet.* Real and live,

Vendor-agnostic. Model-agnostic. I'll apply it on Claude, GPT, Gemini, internal stacks, open-source — whatever fits.

TRAVELS WITH	Any consulting engagement
SELLS AS	Brady's expertise
STATUS	Battle-tested across CPG, QSR, furniture retail

but deliberately not over-publicized. I'd rather build something operators recognize as real before I put it in front of them.

ENTITY	mception LLC
STAGE	Build phase, advisor-supported
STATUS	Not for sale today — applying the framework is the offer

— THE THESIS

Why now, and what I actually believe.

Q If I'm an operator wondering whether to wait, what's your honest answer?

Don't wait. Pick the deepest niche where you can add immediate value if you're better at it in some way — the place where your domain depth is real and your frustrations are sharp. Then start now. Go deep. Start the iterative cycle and let it compound.

Here's why timing matters: **this AI is as dumb as it will ever be.** The next one will be smarter. The one after that, smarter still. When the technology catches up — and in some cases that's a matter of months — the work you did this year doesn't get rebuilt. It gets copy-pasted forward onto a more capable substrate. The frameworks survive. The research compounds. The agents inherit.

The operators who wait for the technology to be "ready" will spend the catch-up window watching the early movers' systems get smarter every month while they're still arguing about pilot scope.

"This AI is as dumb as it will ever be. The frameworks you build today travel forward onto every smarter version that follows."

Q Where does this go in two or three years?

Assuming the major model providers figure out how to interoperate — which is the load-bearing assumption — I expect **~90% of the operating logic runs behind the scenes.** Lower-skill workers maintain the recursive learning process and make sure the rules are still firing correctly. Most rules can be set and enforced at a high level. You stop repeating yourself.

The four-hour workweek becomes plausible — if you define "work" the way most operators currently experience it: sitting at a screen in meetings and answering email. Strip that out, and what's left is the part that was always the actual job: **being in the field, meeting people, exploring, thinking big.** Record it or give a quick debrief afterward, and the system makes it executable.

If anything, the harder problem becomes deciding where to **stop**. The system will do too much on its own if you let it. You'll need to be deliberate about which decisions you keep.

A VIGNETTE · WHERE THIS IS HEADING

If your AI understands your leases better than you do...

...and the language capabilities are good enough to handle counterparties — why wouldn't it negotiate the renewal autonomously, and just escalate the genuinely ambiguous calls to you? The technology is closer than people realize. The interesting question is no longer *can it*. It's *which authority do I delegate, on what terms, and how do I want to be looped in*. That's an operator question, not a technology question. And it's why the framework matters more than the model.

— IN SHORT

Frameworks are how I work. mception is what I'm building. The window is closing fast on the operators who go deep now.

I'm not selling the platform. I'm applying the framework. The companies that pick a deep niche and start the iterative cycle now will spend the next 36 months compounding while everyone else is still scoping pilots. That's the bet I'm making — for myself, for the operators I work with, and for what mception eventually becomes.

Brady Smallwood · Independent strategy and operations consultant focused on AI-native transformation for retail and foodservice operators. Founder of **mception**, an early-stage platform company. Prior roles include COO and board member in multi-unit operating environments. MBA, University of Chicago Booth School of Business.